

Assignment 1, Map and Reduce

Members: Kevin Slachta, Tushar Sharma, Vivek Rajan

Architectural Description:

Our program read input through the command line through arguments. We utilized booleans to determine whether the user is requesting to sort a list or do a word count as well as whether to use processes or threads. Next we took in the file from the user and begin setting up the environment for the mapper process. This involved reading the information into a linked list in order to allow our mapper to evenly divide our data and process it efficiently. Our nodes were structured in a way that we stored the different chunks depending on the number of reduces. Once the file is read in, we forked or created new threads to in such a way that we can perform the word count or the sort on the small chunks of data and then combine the results thereafter.

In order to do the word count we used the same linked list but in this case we generated a key/value pair where key is the unique word and value is the number of times it is seen. Map then puts its list of key/value pairs into shared memory for reduces threads/processes to use. For the sort portion we used a similar architecture and to sort we implemented a merge sort to organize the information. Our reduce function that takes the information from each of the threads/processes in shared memory then combines it into one structured solution. In word count it takes the sum of the same word.

Efficiency:

Through our analysis of thread and processor approach we discovered that a multi-threaded is slightly faster than a multi-processor approach. Our team came to the conclusion that this was due to the fact that when a new process is created the heap must be copied over. It is seen however that this could be a flaw in threads in that one must be careful when using shared memory. Threads that read/write to the same address could lead to inconsistency in results and thus must be handled carefully

Problems:

Our group came across a few difficulties when running creating our algorithm. The biggest problem we came across was a proper way to share counters and other global variables amongst threads. Often

times our code would randomly segfault and the cause was uncertain. Next, we faced an issue with waiting on all our threads to finish. Our parent thread would sometimes finish without all the child threads finishing and our results would be incomplete. Our biggest difficulty with the assignment was a proper way to debug our errors. Because the threads/processors are running simultaneously and independently of one another it became hard to identify when some error occurred. We solved this problem by only working with 2 threads/processes first to verify the correctness of our algorithm and then we were able to scale it to handle more.